

## ARIMA Model and Box-Jenkins Methodology

The ARIMA (AutoRegressive Integrated Moving Average) model is a popular time series forecasting method that captures the relationship between an observation and a number of lagged observations, as well as the lagged forecast errors. The model is defined as:

$$Y_t = \phi_1 Y_{t-1} + \phi_2 Y_{t-2} + \dots + \phi_p Y_{t-p} + \theta_1 \epsilon_{t-1} + \theta_2 \epsilon_{t-2} + \dots + \theta_q \epsilon_{t-q} + \epsilon_t$$

Where:

- $Y_t$  is the actual value at time  $t$ .
- $\phi_1, \phi_2, \dots, \phi_p$  are the autoregressive coefficients.
- $\theta_1, \theta_2, \dots, \theta_q$  are the moving average coefficients.
- $\epsilon_t$  is the error term (white noise).
- $p$  is the order of the autoregressive part.
- $q$  is the order of the moving average part.
- $d$  is the number of differencing required to make the series stationary.

The Box-Jenkins methodology involves the following steps to identify the best ARIMA model:

1. **Model Identification:** Determine the values of  $p$ ,  $d$ , and  $q$  using tools like the Autocorrelation Function (ACF) and Partial Autocorrelation Function (PACF) plots.
2. **Parameter Estimation:** Estimate the coefficients  $\phi$  and  $\theta$  using methods like Maximum Likelihood Estimation (MLE).
3. **Model Checking:** Validate the model by analyzing residuals and ensuring they behave like white noise.
4. **Forecasting:** Use the fitted model to forecast future values.

## Matrices and Equations

For AR(1) process:

$$Y_t = \phi_1 Y_{t-1} + \epsilon_t$$

Matrix form:

$$\begin{pmatrix} Y_t \\ Y_{t-1} \\ \vdots \\ Y_1 \end{pmatrix} = \begin{pmatrix} \phi_1 & 0 & \dots & 0 \\ 1 & 0 & \dots & 0 \\ 0 & 1 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & 1 \end{pmatrix} \begin{pmatrix} Y_{t-1} \\ Y_{t-2} \\ \vdots \\ Y_0 \end{pmatrix} + \begin{pmatrix} \epsilon_t \\ 0 \\ \vdots \\ 0 \end{pmatrix}$$

```

In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import statsmodels.api as sm
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
import yfinance as yf

# Fetch the data from Yahoo Finance
ticker = 'MSFT'
data = yf.download(ticker, start='2010-01-01', end='2024-07-31', interval='1mo')

# Keep only the 'Close' column
data = data['Close']

# Differencing to make the series stationary
data_diff = data.diff().dropna()

# ACF and PACF plots
plt.figure(figsize=(12, 6))
plt.subplot(121)
plot_acf(data_diff, ax=plt.gca(), lags=40)
plt.title('Autocorrelation Function (ACF)')

plt.subplot(122)
plot_pacf(data_diff, ax=plt.gca(), lags=40)
plt.title('Partial Autocorrelation Function (PACF)')
plt.show()

# Model Identification: Using AIC/BIC to select the best ARIMA model
best_aic = np.inf
best_order = None
best_model = None

for p in range(5):
    for d in range(2):
        for q in range(5):
            try:
                model = ARIMA(data, order=(p, d, q)).fit()
                if model.aic < best_aic:
                    best_aic = model.aic
                    best_order = (p, d, q)
                    best_model = model
            except Exception as e:
                print(f"ARIMA({p},{d},{q}) failed to fit: {e}")
                continue

if best_model is not None:
    print(f'Best ARIMA model: {best_order} with AIC: {best_aic}')
else:
    print("No suitable ARIMA model found.")

# Display statistical results of the best ARIMA model
if best_model is not None:
    print(best_model.summary())

# Plot AR and MA roots in the unit circle
if best_model is not None and best_order[0] > 0: # Only plot if AR or MA parameters
    ar_params = np.r_[1, -best_model.arparams] if best_order[0] > 0 else [1]
    ma_params = np.r_[1, best_model.maparams] if best_order[2] > 0 else [1]

    ar_roots = np.roots(ar_params)

```

```

ma_roots = np.roots(ma_params)

plt.figure(figsize=(8, 8))
unit_circle = plt.Circle((0, 0), 1, color='blue', fill=False, linestyle='--')
plt.gca().add_patch(unit_circle)

plt.scatter(ar_roots.real, ar_roots.imag, color='red', label='AR Roots')
plt.scatter(ma_roots.real, ma_roots.imag, color='green', label='MA Roots')

plt.axhline(0, color='black', linewidth=0.5)
plt.axvline(0, color='black', linewidth=0.5)
plt.xlim(-2, 2)
plt.ylim(-2, 2)
plt.gca().set_aspect('equal', adjustable='box')

plt.title('AR and MA Roots in the Unit Circle')
plt.legend()
plt.show()

# Unit Root Graph (ADF Test)
adf_test = sm.tsa.adfuller(data_diff)
print(f'ADF Statistic: {adf_test[0]}')
print(f'p-value: {adf_test[1]}')

# Residuals Plot
if best_model is not None:
    residuals = pd.DataFrame(best_model.resid)
    residuals.plot(title="Residuals")
    plt.show()

# Actual vs Forecast
plt.figure(figsize=(10, 6))
plt.plot(data, label='Actual')
plt.plot(best_model.fittedvalues, color='red', label='Fitted')
plt.legend()
plt.title('Actual vs Fitted')
plt.show()


# Forecasting until the end of 2025
forecast_steps = 17 # Number of months from August 2024 to December 2025
forecast = best_model.get_forecast(steps=forecast_steps)
forecast_index = pd.date_range(start=data.index[-1], periods=forecast_steps+1, f

# Plot Actual vs Forecast
plt.figure(figsize=(10, 6))
plt.plot(data, label='Actual')
plt.plot(best_model.fittedvalues, color='red', label='Fitted')
plt.plot(forecast_index, forecast.predicted_mean, color='green', label='Forecast')
plt.fill_between(forecast_index, forecast.conf_int()[0], forecast.conf_int()[1])
plt.legend()
plt.title('Actual vs Fitted and Forecasted Values')
plt.show()

else:
    print("Skipping residual and forecast plots since no model was fitted.")

```

[\*\*\*\*\*100%\*\*\*\*\*] 1 of 1 completed



C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:581: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

warnings.warn('A date index has been provided, but it has no')

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:581: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

warnings.warn('A date index has been provided, but it has no')

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:581: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

warnings.warn('A date index has been provided, but it has no')

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:581: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

warnings.warn('A date index has been provided, but it has no')

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:581: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

warnings.warn('A date index has been provided, but it has no')

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:581: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

warnings.warn('A date index has been provided, but it has no')

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\statespace\sarimax.py:978: UserWarning: Non-invertible starting MA parameters found. Using zeros as starting parameters.

warn('Non-invertible starting MA parameters found.')

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:581: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

warnings.warn('A date index has been provided, but it has no')

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:581: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

warnings.warn('A date index has been provided, but it has no')

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:581: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

warnings.warn('A date index has been provided, but it has no')

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:581: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

warnings.warn('A date index has been provided, but it has no')

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:581: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

warnings.warn('A date index has been provided, but it has no')

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:581: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

warnings.warn('A date index has been provided, but it has no')

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:581: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

warnings.warn('A date index has been provided, but it has no')

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\base\model.py:566: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle\_retvals

warnings.warn("Maximum Likelihood optimization failed to ")

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:581: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

warnings.warn('A date index has been provided, but it has no')

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:581: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

warnings.warn('A date index has been provided, but it has no')

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:581: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

warnings.warn('A date index has been provided, but it has no')

C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\base\model.py:566: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle\_retvals

warnings.warn("Maximum Likelihood optimization failed to ")

Best ARIMA model: (2, 1, 1) with AIC: 1346.4776082955477

#### SARIMAX Results

```
=====
Dep. Variable:          Close    No. Observations:          233
Model:                ARIMA(2, 1, 1)    Log Likelihood          -669.239
Date:                Thu, 08 Aug 2024    AIC                  1346.478
Time:                21:59:27    BIC                  1360.265
Sample:                0    HQIC                  1352.038
                        - 233
```

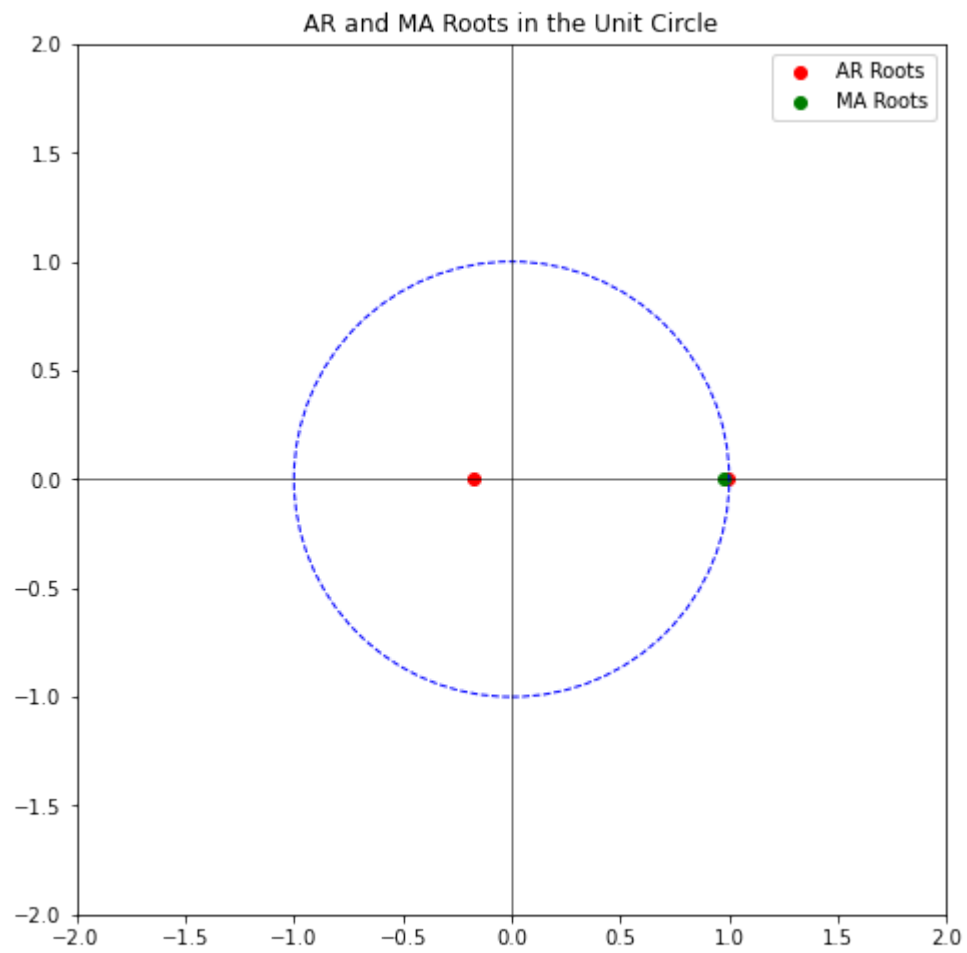
Covariance Type: opg

```
=====
              coef    std err          z      P>|z|      [0.025      0.975]
-----
ar.L1         0.8181     0.054    15.068     0.000     0.712     0.925
ar.L2         0.1778     0.053     3.358     0.001     0.074     0.282
ma.L1        -0.9783     0.030   -32.550     0.000    -1.037    -0.919
sigma2       106.2550     6.499    16.349     0.000    93.517   118.993
=====
```

```
=====
Ljung-Box (L1) (Q):                0.09    Jarque-Bera (JB):                488.43
Prob(Q):                          0.76    Prob(JB):                      0.00
Heteroskedasticity (H):            106.62    Skew:                          0.42
Prob(H) (two-sided):              0.00    Kurtosis:                      10.06
=====
```

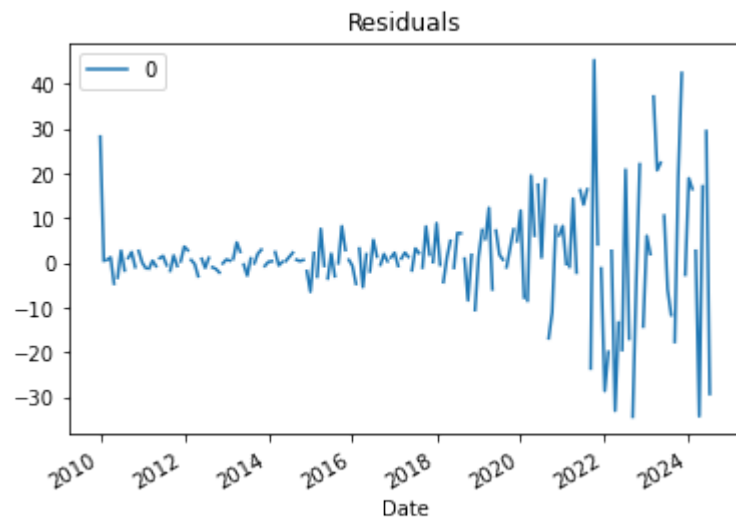
#### Warnings:


[1] Covariance matrix calculated using the outer product of gradients (complex-step).



ADF Statistic: -9.79453953047507

p-value: 6.224781218168595e-17





C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\tsa\base\tsa\_model.py:376: ValueWarning: No supported index is available. Prediction results will be given with an integer index beginning at `start`.

warnings.warn('No supported index is available.')

```
-----
TypeError                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_9456\464922073.py in <module>
    109     plt.plot(best_model.fittedvalues, color='red', label='Fitted')
    110     plt.plot(forecast_index, forecast.predicted_mean, color='green', label='F
forecast')
--> 111     plt.fill_between(forecast_index, forecast.conf_int()[0], forecast.conf
_int()[1], color='green', alpha=0.2)
    112     plt.legend()
    113     plt.title('Actual vs Fitted and Forecasted Values')
```

```
C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\frame.py in __getitem__(self,
key)
```

```
    3456         if self.columns.nlevels > 1:
    3457             return self._getitem_multilevel(key)
-> 3458         indexer = self.columns.get_loc(key)
    3459         if is_integer(indexer):
    3460             indexer = [indexer]
```


```
C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\indexes\base.py in get_loc(self,
key, method, tolerance)
```

```
    3359         casted_key = self._maybe_cast_indexer(key)
    3360         try:
-> 3361             return self._engine.get_loc(casted_key)
    3362         except KeyError as err:
    3363             raise KeyError(key) from err
```

```
C:\ProgramData\Anaconda3\lib\site-packages\pandas\_libs\index.pyx in pandas._libs.ind
ex.IndexEngine.get_loc()
```

```
C:\ProgramData\Anaconda3\lib\site-packages\pandas\_libs\index.pyx in pandas._libs.ind
ex.IndexEngine.get_loc()
```

**TypeError:** '(slice(None, None, None), 0)' is an invalid key



In [ ]: